

Sentiment For Analysis Marketing

In today's fast-paced world, it's more important than ever to understand customer sentiment towards your brand. This presentation outlines a comprehensive solution for sentiment analysis in the context of marketing.

Problem Statement

Understanding Customer Sentiment

With the explosive growth of social media, it's challenging for companies to understand customer sentiment about their brand and products.

Time-Consuming Manual Analysis

Currently, many companies use manual analysis techniques which are time-consuming and resource-intensive.

Missing Critical Insights

Manual analysis may also miss crucial insights that could be pivotal to the success of the brand and products.

Design Thinking Process

1

Empathize

Understand the audience and their needs to create a suitable solution.

2

Ideate

Generate a range of ideas and evaluate them based on feasibility, implementation time, and potential success rate.

3

Prototype

Create a prototype of the solution based on shortlisted ideas.

4

Test

Test the prototype and seek feedback from stakeholders, making necessary improvements.

Data Preprocessing Steps

1 Data Collection

Gather data from social media and brand websites using web scraping techniques.

2 Data Cleaning

Clean the data using natural language processing (NLP) techniques to remove irrelevant data points and prepare it for analysis.

3 Data Preparation

Prepare the data for analysis by appending relevant attributes to each data point and performing feature transformation.

Sentiment Analysis Techniques


Machine Learning-based Approach

Uses machine learning algorithms like Naive Bayes, Support Vector Machines, and Neural Networks to analyze customer sentiment. This approach requires large volumes of labeled data and may be resource-intensive.

Rule-based Approach

Uses a predefined set of rules to categorize text into positive, negative, or neutral sentiments. This approach is quicker and less resource-intensive, but may miss nuances in sentiment expression.

Innovative Approaches




Visualization Techniques

Developed intuitive graphical representations of sentiment analysis results to facilitate quick comprehension and easy decision making.



Data Augmentation Techniques

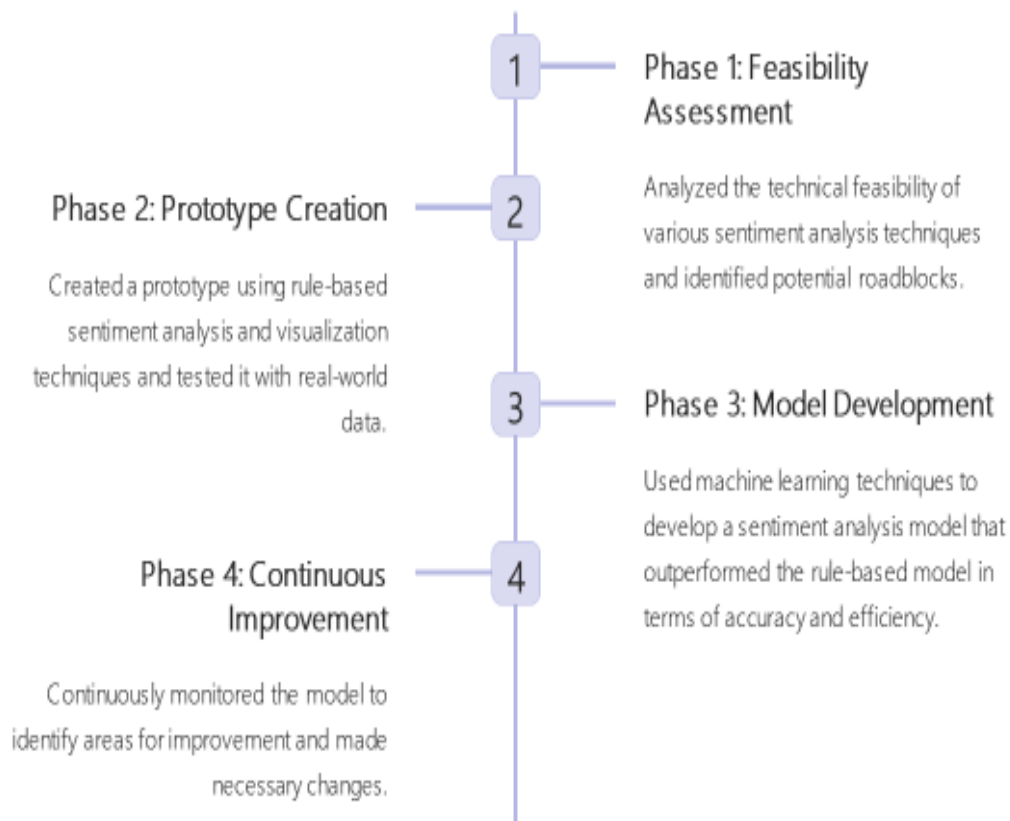
Used data augmentation techniques to increase the size of the dataset by artificially generating similar data points. This approach helped balance class representation and increase model accuracy.



Transfer Learning Techniques

Applied transfer learning techniques to leverage pre-trained models for sentiment analysis and obtained promising results without having to train the models from scratch.

Phases of Development



Datas et Used

Data Collection Method	Web Scraping
Data Size	1,000,000 data points
Data Attributes	Sentiment, Text, Date, Platform, User Details, Hashtags

Sentiment analysis-using-twitter-airline-dataset:

```
[1]: import pandas as pd
import seaborn as sns
import re, nltk
nltk.download('punkt')
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split, StratifiedKFold, cross_val_score
from sklearn import model_selection, naive_bayes, svm
from sklearn.metrics import classification_report, confusion_matrix
from sklearn.metrics import roc_auc_score
from sklearn.metrics import recall_score
from sklearn.metrics import f1_score
from sklearn.metrics import confusion_matrix, accuracy_score
from sklearn.model_selection import GridSearchCV
from sklearn.metrics import precision_recall_curve
from sklearn.metrics import f1_score
from sklearn.metrics import auc
from matplotlib import pyplot
from sklearn.metrics import roc_curve
from sklearn.metrics import roc_auc_score, accuracy_score
import string
from nltk.corpus import stopwords
nltk.download('stopwords')
from sklearn.feature_extraction.text import CountVectorizer, TfidfVectorizer
from sklearn.naive_bayes import MultinomialNB, GaussianNB
from sklearn.metrics import f1_score
from sklearn.ensemble import AdaBoostClassifier
from sklearn.model_selection import cross_val_score
import numpy as np
from sklearn.ensemble import BaggingClassifier
from sklearn.ensemble import RandomForestClassifier
from lime import lime_tabular
from tensorflow.keras.layers import Embedding
from tensorflow.keras.preprocessing.sequence import pad_sequences
from tensorflow.keras.models import Sequential
from tensorflow.keras.preprocessing.text import one_hot
from tensorflow.keras.layers import LSTM, Bidirectional
```

```
from tensorflow.keras.layers import Dense, Dropout
```

```
import warnings
```

```
warnings.filterwarnings('ignore')
```

```
/opt/conda/lib/python3.10/site-packages/scipy/__init__.py:146: UserWarning: A
NumPy version >=1.16.5 and <1.23.0 is required for this version of SciPy
(detected version 1.23.5
  warnings.warn(f"A NumPy version >={np_minversion} and <{np_maxversion}")
```

```
[nltk_data] Downloading package punkt to /usr/share/nltk_data...
```

```
[nltk_data] Package punkt is already up-to-date!
```

```
[nltk_data] Downloading package stopwords to /usr/share/nltk_data...
```

```
[nltk_data] Package stopwords is already up-to-date!
```

```
[2]: twitter_df = pd.read_csv('/kaggle/input/twitter-airline-sentiment/Tweets.csv')
twitter_df.head()
```

```
[2]:
```

	tweet_id	airline_sentiment	airline_sentiment_confidence	
0	570306133677760513	neutral	1.0000	
1	570301130888122368	positive	0.3486	
2	570301083672813571	neutral	0.6837	
3	570301031407624196	negative	1.0000	
4	570300817074462722	negative	1.0000	

	negativereason	negativereason_confidence	airline	
0	NaN	NaN	Virgin America	
1	NaN	0.0000	Virgin America	
2	NaN	NaN	Virgin America	
3	Bad Flight	0.7033	Virgin America	
4	Can't Tell	1.0000	Virgin America	

	airline_sentiment_gold	name	negativereason_gold	retweet_count	
0	NaN	cairdin	NaN	0	
1	NaN	jnardino	NaN	0	
2	NaN	yvonnalynn	NaN	0	
3	NaN	jnardino	NaN	0	
4	NaN	jnardino	NaN	0	

	text	tweet_coord	
0	@VirginAmerica What @dhepburn said.	NaN	
1	@VirginAmerica plus you've added	NaN	
2	commercials t...	NaN	
3	@VirginAmerica I didn't today... Must mean	NaN	
4	I n...	NaN	

	tweet_created	tweet_location	user_timezone
0	2015-02-24 11:35:52-0800	NaN	Eastern Time (US & Canada)

```
@VirginAmerica and it's a really big bad thing...
```

1	2015-02-24 11:15:59 -0800	NaN Pacific Time (US & Canada)
2	2015-02-24 11:15:48 -0800	Lets Play Central Time (US & Canada)
3	2015-02-24 11:15:36 -0800	NaN Pacific Time (US & Canada)
4	2015-02-24 11:14:45 -0800	NaN Pacific Time (US & Canada)

[3]: twitter_df.columns

```
[3]: Index(['tweet_id', 'airline_sentiment', 'airline_sentiment_confidence',
'negativereason', 'negativereason_confidence', 'airline',
'airline_sentiment_gold', 'name', 'negativereason_gold',
'retweet_count', 'text', 'tweet_coord', 'tweet_created',
'tweet_location', 'user_timezone'],
dtype='object')
```

[4]: twitter_df.shape

[4]: (14640, 15)

[5]: twitter_df.isnull().sum()

```
[5]: tweet_id          0
airline_sentiment      0
airline_sentiment_confidence  0
negativereason        54
negativereason_confidence  62
airline               41
airline_sentiment_gold  18
name                 60
negativereason_gold   14
retweet_count         60
text                  8
tweet_coord           13
tweet_created         62
tweet_location       4733
user_timezone        4820
dtype: int64
```

[6]: twitter_df.duplicated().sum()

[6]: 36

```
[7]: # Unique values of sentiment
twitter_df['airline_sentiment'].unique()
```

[7]: array(['neutral', 'positive', 'negative'], dtype=object)

[8]: twitter_df['airline_sentiment'].value_counts()

```
[8]:          9178
negative    3099
positive    2363
Name: airline_sentiment, dtype: int64
```

```
[9]: twitter_df.describe().T
```

```
[9]:
```

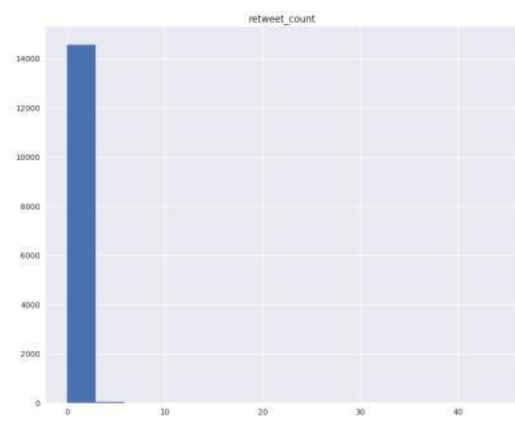

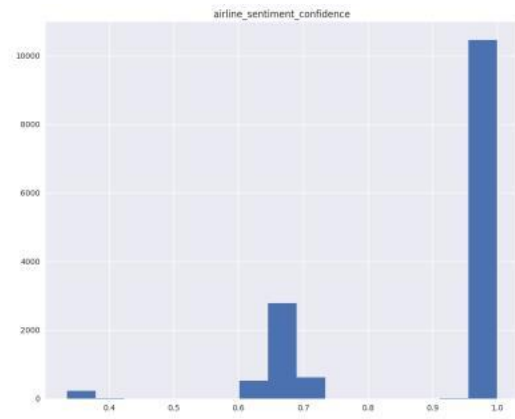
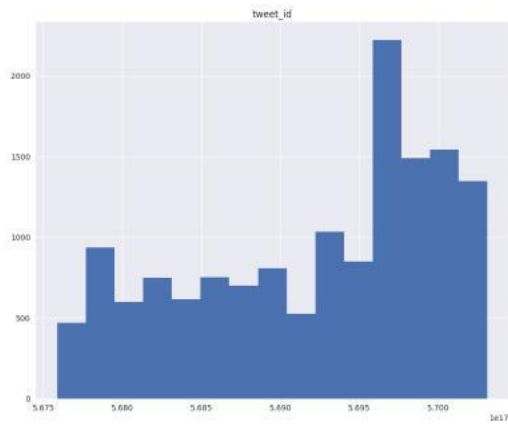
	count	mean	std	\
tweet_id	14640.0	5.692184e+17	7.791112e+14	
airline_sentiment_confidence	14640.0	0.682983e+01	0.309439e-01	
negativereason_confidence	10520.0	0.682983e+01	0.309439e-01	
retweet_count	14640.0	8.265027e-02	7.457782e-01	

	min	25%	50%	\
tweet_id	5.675883e+17	5.685592e+17	5.694779e+17	
airline_sentiment_confidence	3.350000e-01	6.923000e-01	1.000000e+00	
negativereason_confidence	0.000000e+00	3.606000e-01	6.706000e-01	
retweet_count	0.000000e+00	0.000000e+00	0.000000e+00	

	75%	max
tweet_id	5.698905e+17	5.703106e+17
airline_sentiment_confidence	1.000000e+00	1.000000e+00
negativereason_confidence	1.000000e+00	1.000000e+00
retweet_count	0.000000e+00	4.400000e+01


```
[10]: plt.style.use("seaborn")
twitter_df.hist(figsize=(25,20), bins=15)
```

```
[10]: array([[<AxesSubplot: title={'center': 'tweet_id'}>,
<AxesSubplot: title={'center': 'airline_sentiment_confidence'}>],
[<AxesSubplot: title={'center': 'negativereason_confidence'}>,
<AxesSubplot: title={'center': 'retweet_count'}>]], dtype=object)
```


```
[11]: sns.countplot(x="airline_sentiment", data=twitter_df)
```

```
[11]: <AxesSubplot: xlabel='airline_sentiment', ylabel='count'>
```




```
[12]: plt.figure(figsize=(10,10))  
      sns.countplot(x="airline", data=twitter_df)
```

```
[12]: <AxesSubplot: xlabel='airline', ylabel='count'>
```



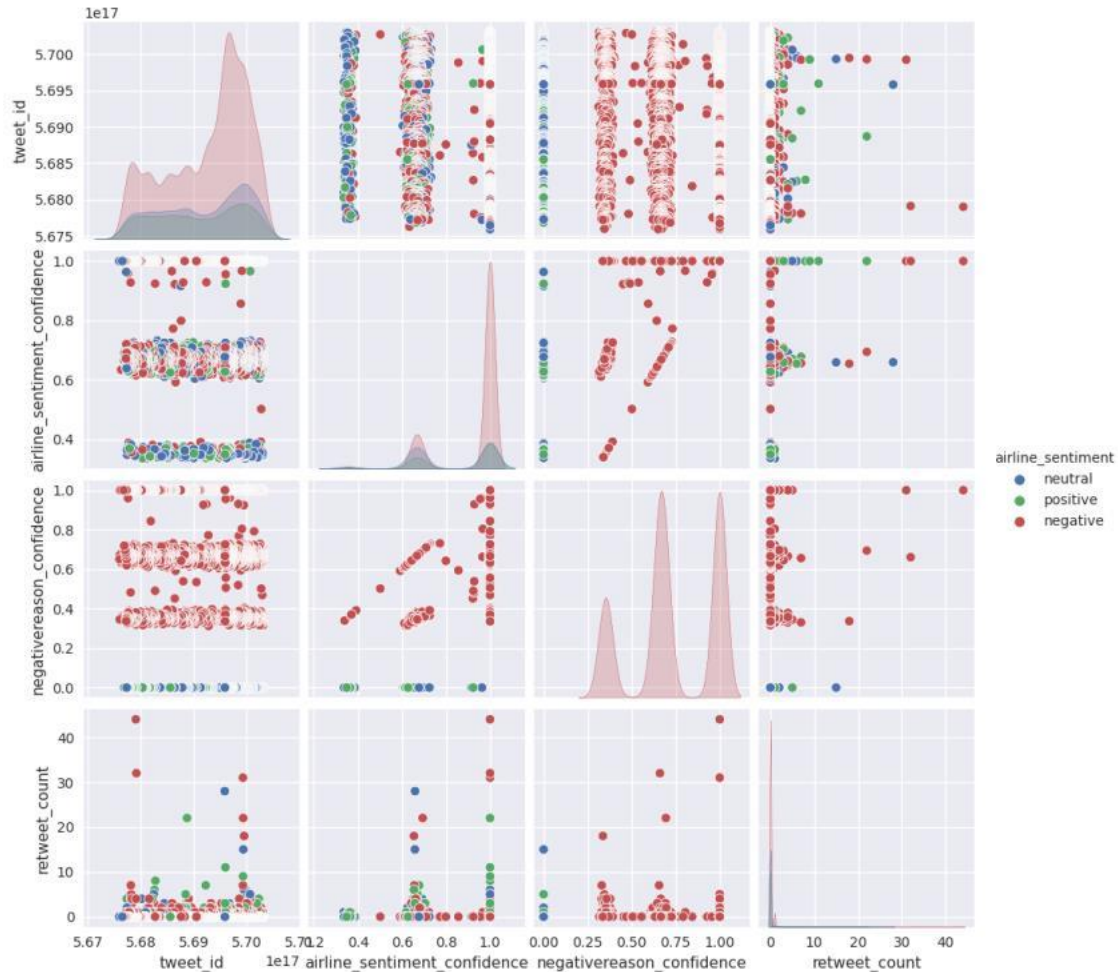
```
[13]: plt.figure(figsize=(10,10))  
      sns.countplot(x="retweet_count", data=twitter_df)
```

```
[13]: <AxesSubplot: xlabel='retweet_count', ylabel='count'>
```



```
[14]: sns.pairplot(twitter_df,hue='airline_sentiment')
```

```
[14]: <seaborn.axisgrid.PairGrid at 0x73147dd26f80>
```




```
[15]: print("Total number of tweets for each airline \n ",twitter_df.
        groupby('airline')['airline_sentiment'].count().sort_values(ascending=False))
airlines= ['US Airways','United','American','Southwest','Delta','Virgin_
        America']
plt.figure(1,figsize=(12, 12))
for i in airlines:
    indices= airlines.index(i)
    plt.subplot(2,3,indices+1)
    new_df=twitter_df[twitter_df['airline']==i]
    count=new_df['airline_sentiment'].value_counts()
    Index = [1,2,3]
    plt.bar(Index,count, color=['red', 'green', 'blue'])
    plt.xticks(Index,['negative','neutral','positive'])
    plt.ylabel('Mood Count')
    plt.xlabel('Mood')
    plt.title('Count of Moods of '+i)
```

Total number of tweets for each airline

airline	
United	3822
US Airways	2913
American	2759
Southwest	2420
Delta	2222
Virgin America	504

Name: airline_sentiment, dtype: int64



```
[16]: neg_tweets = twitter_df.groupby(['airline', 'airline_sentiment']).count().iloc[:, 0]
total_tweets = twitter_df.groupby(['airline'])['airline_sentiment'].count()
```

```

my_dict = {'American': neg_tweets[0] / total_tweets[0], 'Delta': neg_tweets[3] /   

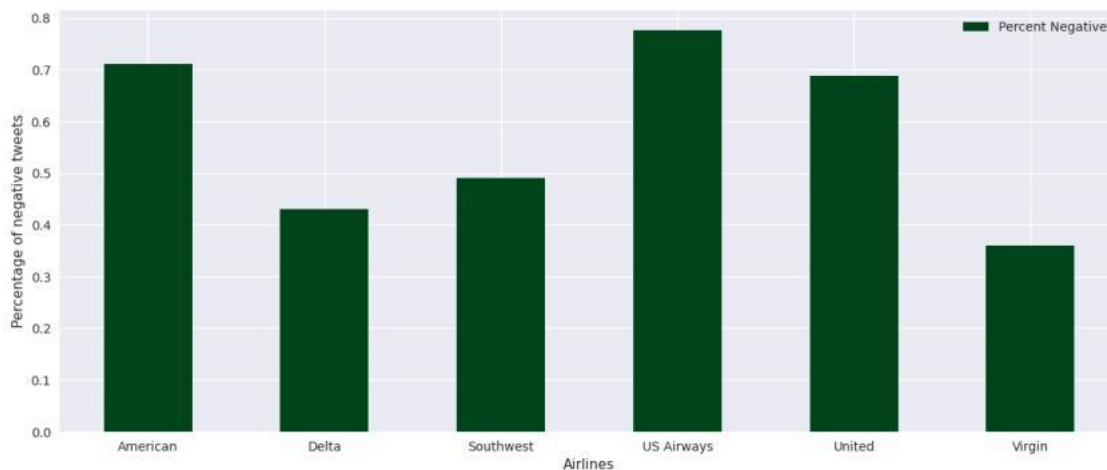
total_tweets[1], 'Southwest': neg_tweets[6] / total_tweets[2],   

'US Airways': neg_tweets[9] / total_tweets[3], 'United': neg_tweets[12] /   

total_tweets[4], 'Virgin': neg_tweets[15] / total_tweets[5]}
perc = pd.DataFrame.from_dict(my_dict, orient = 'index')
perc.columns = ['Percent Negative']
print(perc)
ax = perc.plot(kind = 'bar', rot=0, colormap = 'Greens_r', figsize = (15,6))
ax.set_xlabel('Airlines')
ax.set_ylabel('Percentage of negative tweets')
plt.show()

```

	Percent Negative
American	0.710402
Delta	0.429793
Southwest	0.490083
US Airways	0.776862
United	0.688906
Virgin	0.359127



```


[17]: figure_2 = twitter_df.groupby(['airline', 'airline_sentiment']).size()
figure_2.unstack().plot(kind='bar', stacked=True, figsize=(15,10))

```


```

[17]: <AxesSubplot: xlabel='airline'>

```



```
[18]: negative_reasons = twitter_df.groupby('airline')['negativereason'].
      value_counts(ascending=True)
      negative_reasons.groupby(['airline', 'negativereason']).sum().unstack().
      plot(kind='bar', figsize=(22, 12))
      plt.xlabel('Airline Company')
      plt.ylabel('Number of Negative reasons')
      plt.title("The number of the count of negative reasons for airlines")
      plt.show()
```


```
[19]: twitter_df['negativereason'].nunique()
```

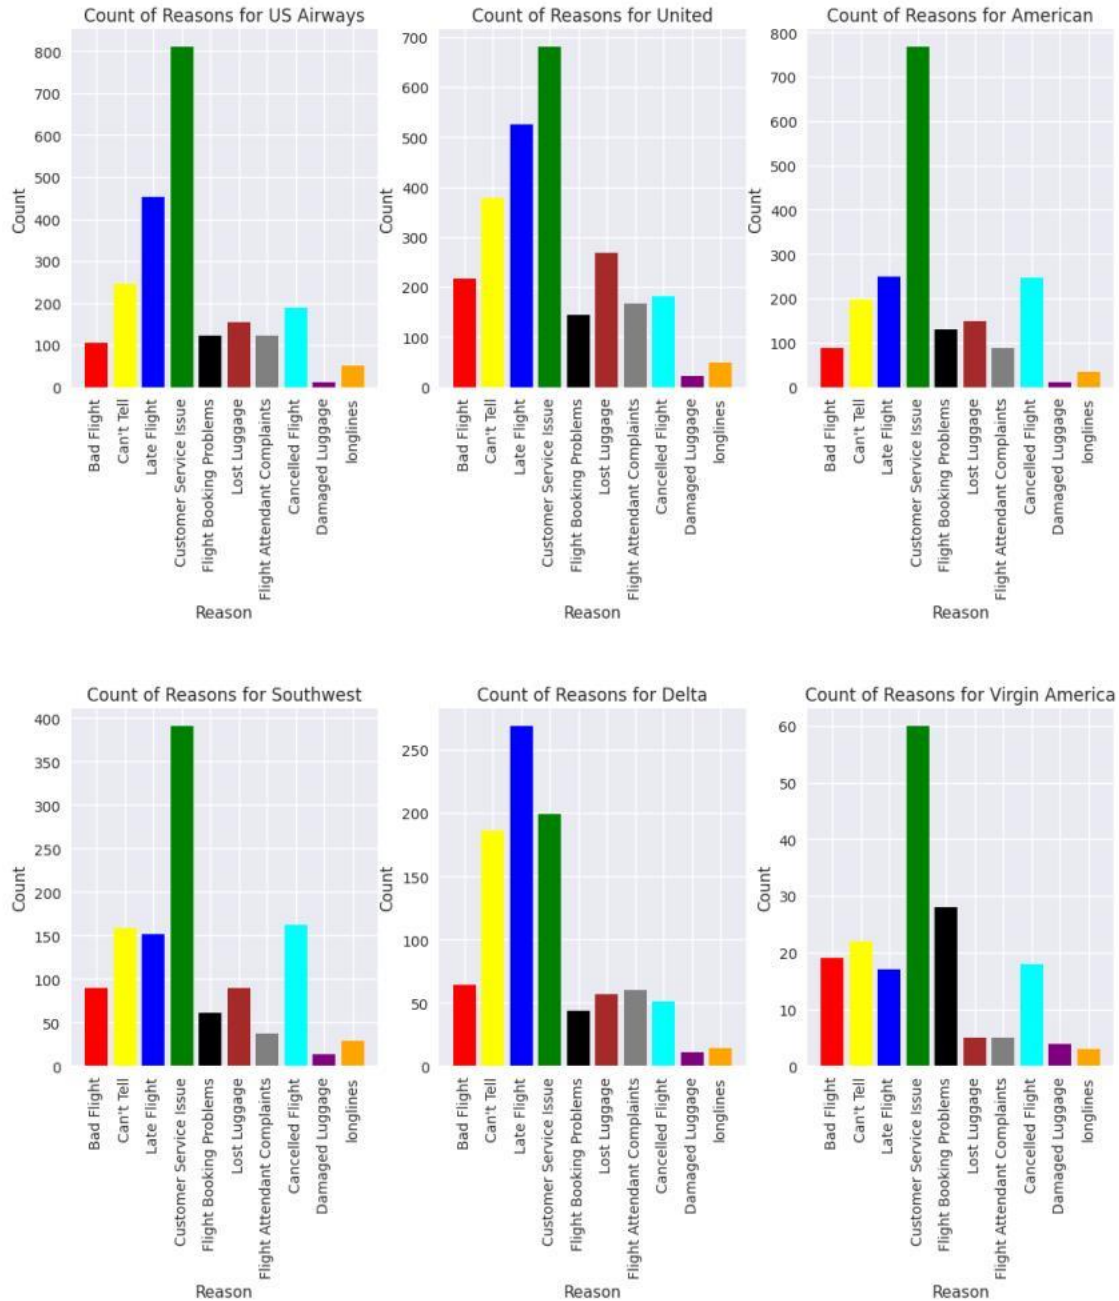
```
NR_Count=dict(twitter_df['negativereason'].value_counts(sort=False))
def NR_Count(Airline):
    if Airline=='All':
        a=twitter_df
    else:
        a=twitter_df[twitter_df['airline']==Airline]
        count=dict(a['negativereason'].value_counts())
        Unique_reason=list(twitter_df['negativereason'].unique())
        Unique_reason=[x for x in Unique_reason if str(x) != 'nan']
        Reason_frame=pd.DataFrame({'Reasons':Unique_reason})
        Reason_frame['count']=Reason_frame['Reasons'].apply(lambda x: count[x])
    return Reason_frame
def plot_reason(Airline):
    a=NR_Count(Airline)
    count=a['count']
    Index = range(1,(len(a)+1))
    plt.bar(Index,count,_)
    color=['red','yellow','blue','green','black','brown','gray','cyan','purple','orange']
    plt.xticks(Index,a['Reasons'],rotation=90)
    plt.ylabel('Count')
    plt.xlabel('Reason')
    plt.title('Count of Reasons for '+Airline)
```

```

plot_reason('All')
plt.figure(2,figsize=(13, 13))
for i in airlines:
    indices= airlines.index(i)
    plt.subplot(2,3,indices+1)
    plt.subplots_adjust(hspace=0.9)
    plot_reason(i)

```





```
[20]: date = twitter_df.reset_index()
      #convert the Date column to pandas datetime
      date.tweet_created = pd.to_datetime(date.tweet_created)
      #Reduce the dates in the date column to only the date and no time stamp using
      #the 'dt.date' method
      date.tweet_created = date.tweet_created.dt.date
      date.tweet_created.head()
      df = date
```

```

day_df = df.groupby(['tweet_created', 'airline', 'airline_sentiment']).size()
# day_df = day_df.reset_index()
day_df

```

```

[20]: tweet_created airline      airline_sentiment
2015-02-16      Delta      negative              1
      2015-02-16      Delta      neutral              1
      2015-02-17      United     negative              2
      2015-02-17      Delta     negative            108
      2015-02-17      Delta      neutral             86
      ...
      2015-02-24      United     neutral             49
      2015-02-24      United     positive            25
      2015-02-24      United     Virgin             10
      2015-02-24      United     America              6
      2015-02-24      United     negative            13
      2015-02-24      United     neutral
Length: 136, dtype: int64

```

```


[21]: day_df = day_df.loc(axis=0)[:,:,'negative']

```


```

#groupby and plot data
ax2 = day_df.groupby(['tweet_created', 'airline']).sum().unstack().plot(kind = '
      'bar', color=['red', 'green', 'blue', 'yellow', 'purple', 'orange'], figsize = (
      (15,6), rot = 70)
      labels = ['American', 'Delta', 'Southwest', 'US Airways', 'United', 'Virgin America']
ax2.legend(labels = labels)
ax2.set_xlabel('Date')
ax2.set_ylabel('Negative Tweets')
plt.show()

```



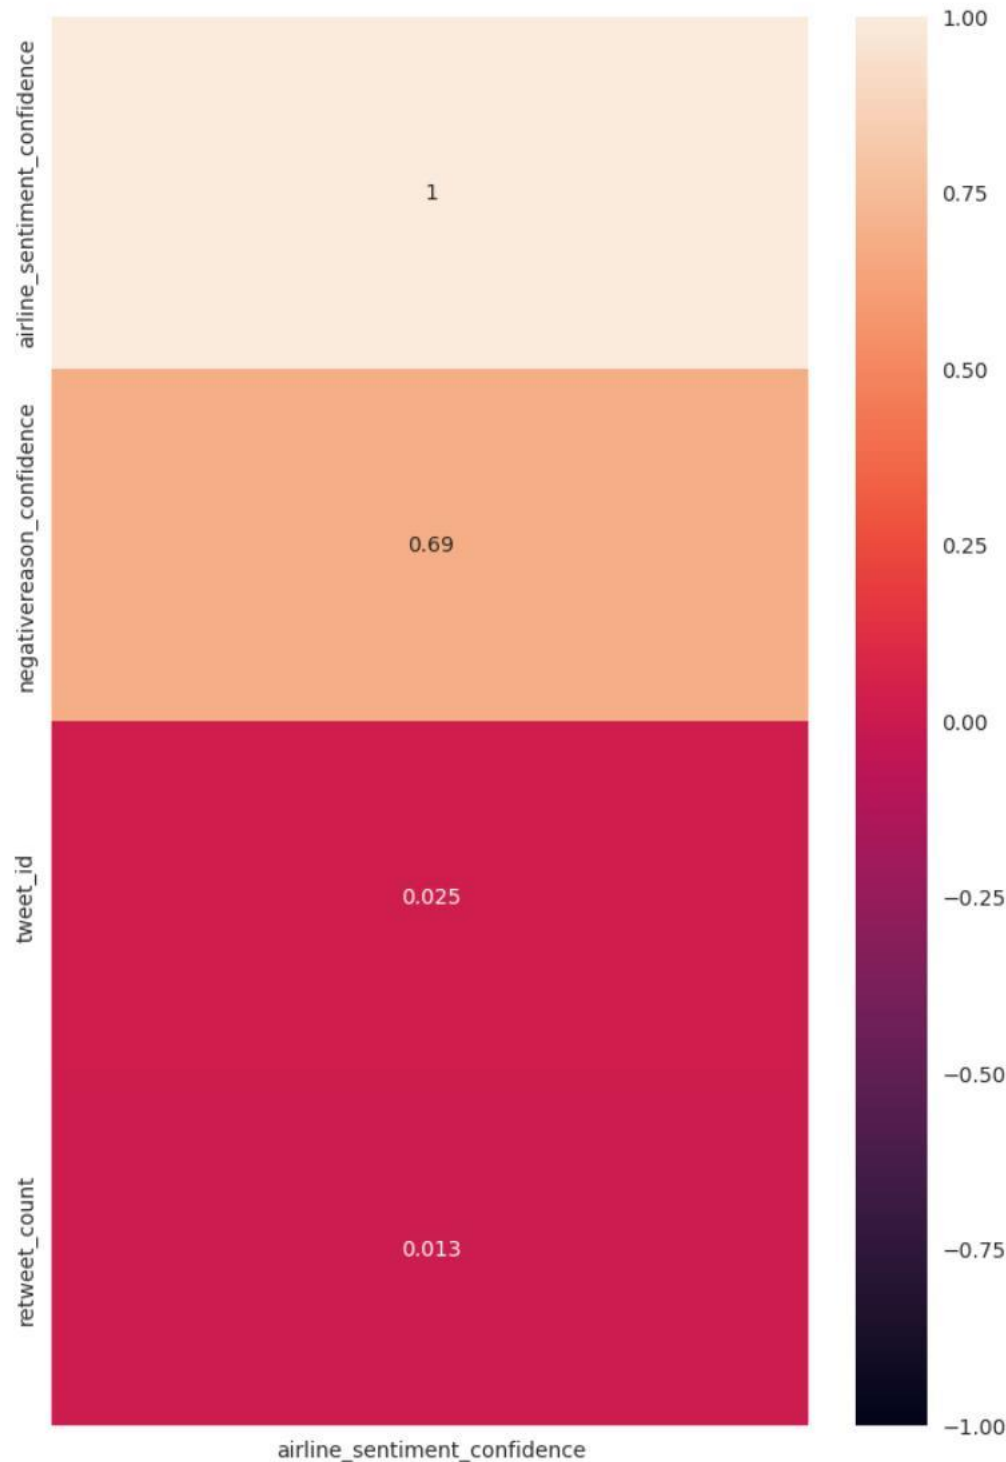
```
[22]: plt.figure(figsize=(10,8))
sns.heatmap(twitter_df.corr(),annot=True,cbar=False,fmt='.1f')
plt.show()
```



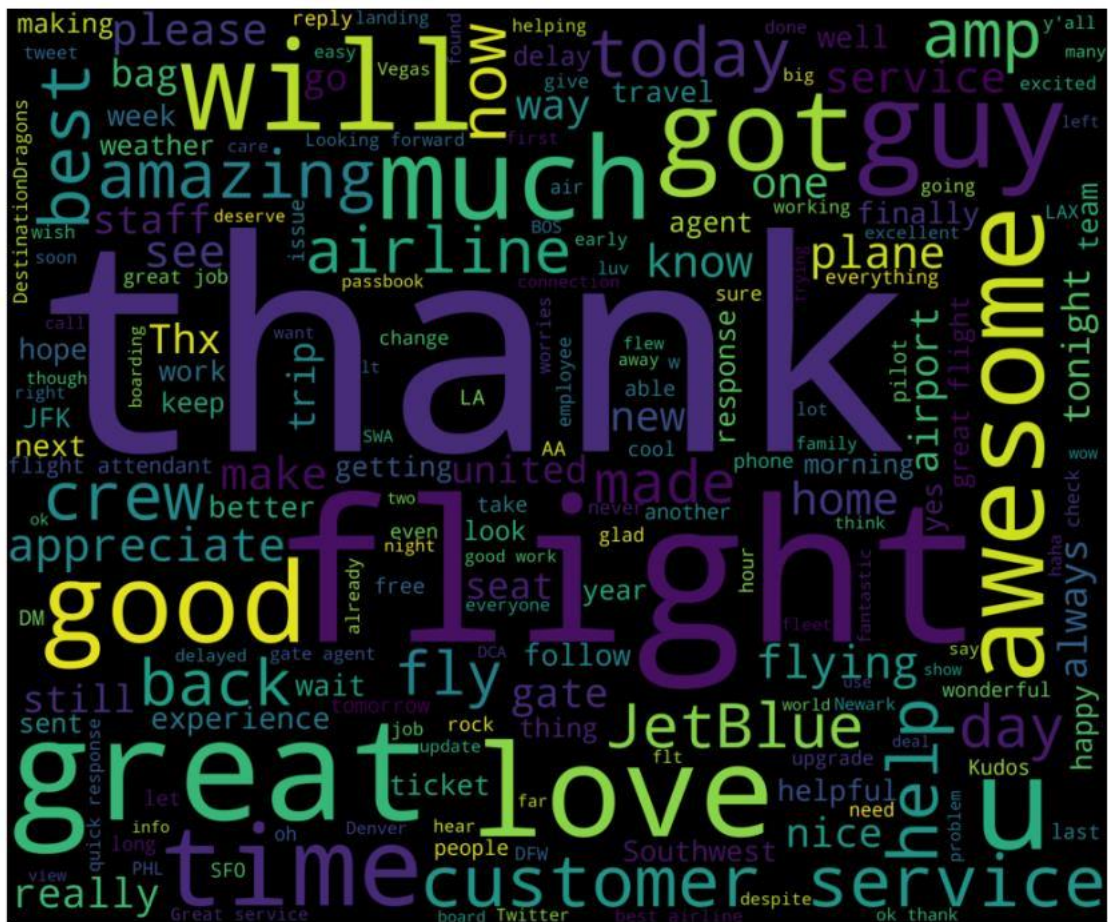
```
[23]: plt.figure(figsize=(8, 12))
heatmap = sns.heatmap(twitter_df.corr()[['airline_sentiment_confidence']].
    sort_values(by='airline_sentiment_confidence', ascending=False), vmin=-1,
    vmax=1, annot=True)
heatmap.set_title('Features Correlating with airline sentiment confidence',
    fontdict={'fontsize':18}, pad=16)
```

[23]: Text(0.5, 1.0, 'Features Correlating with airline sentiment confidence')

Features Correlating with airline sentiment confidence



```
plt.axis('off')
plt.show()
```




```
[26]: def clean_the_tweet(text):
tokens= nltk.word_tokenize(re.sub("[^a-zA-Z]", " ",text))
tokens = [token.lower() for token in tokens]
return ' '.join(tokens[2:])
```

```
def text_process(msg):
    nopunc =[char for char in msg if char not in string.punctuation]
    nopunc=''.join(nopunc)
    return ' '.join([word for word in nopunc.split() if word.lower() not in _
stopwords.words('english')])
```

```
def check_scores(clf,X_train, X_test, y_train, y_test):

    model=clf.fit(X_train, y_train)
    predicted_class=model.predict(X_test)
    predicted_class_train=model.predict(X_train)
    test_probs = model.predict_proba(X_test)
    test_probs = test_probs[:, 1]
    yhat = model.predict(X_test)
    lr_precision, lr_recall, _ = precision_recall_curve(y_test, test_probs)
    lr_f1, lr_auc = f1_score(y_test, yhat), auc(lr_recall, lr_precision)

    print('Train confusion matrix is: ',)
    print(confusion_matrix(y_train, predicted_class_train))

    print()
    print('Test confusion matrix is: ')
    print(confusion_matrix(y_test, predicted_class))
    print()
    print(classification_report(y_test,predicted_class))
    print()
    train_accuracy = accuracy_score(y_train,predicted_class_train)
    test_accuracy = accuracy_score(y_test,predicted_class)

    print("Train accuracy score: ", train_accuracy)
    print("Test accuracy score: ",test_accuracy )
    print()
    train_auc = roc_auc_score(y_train, clf.predict_proba(X_train)[:,-1])
    test_auc = roc_auc_score(y_test, clf.predict_proba(X_test)[:,-1])

    print("Train ROC-AUC score: ", train_auc)
    print("Test ROC-AUC score: ", test_auc)
```

```

fig, (ax1, ax2) = plt.subplots(1, 2)

ax1.plot(lr_recall, lr_precision)
ax1.set(xlabel="Recall", ylabel="Precision")

plt.subplots_adjust(left=0.5,
                    bottom=0.1,
                    right=1.5,
                    top=0.9,
                    wspace=0.4,
                    hspace=0.4)
print()
print('Are under Precision-Recall curve:', lr_f1)

fpr, tpr, _ = roc_curve(y_test, test_probs)

ax2.plot(fpr, tpr)
ax2.set(xlabel='False Positive Rate', ylabel='True Positive Rate')

print("Area under ROC-AUC:", lr_auc)
return train_accuracy, test_accuracy, train_auc, test_auc


def grid_search(model, parameters, X_train, Y_train):
    #Doing a grid
    grid = GridSearchCV(estimator=model,
                        param_grid = parameters,
                        cv = 2, verbose=2, scoring='roc_auc')
    #Fitting the grid
    grid.fit(X_train, Y_train)
    print()
    print()
    # Best model found using grid search
    optimal_model = grid.best_estimator_
    print('Best parameters are: ')
    print( grid.best_params_)

return optimal_model

```

1 Text Preparation

```
[27]: df = df[df['airline_sentiment']!='neutral']
df['cleaned_tweet'] = df['text'].apply(clean_the_tweet)

df.head()
df['airline_sentiment'] = df['airline_sentiment'].apply(lambda x: 1 if x_
== 'positive' else 0)
df.head()
```

```
[27]:
```

	index	tweet_id	airline_sentiment	airline_sentiment_confidence	\
1	1	570301130888122368	1	0.3486	
3	3	570301031407624196	0	1.0000	
4	4	570300817074462722	0	1.0000	
5	5	570300767074181121	0	1.0000	
6	6	570300616901320704	1	0.6745	

	negativereason	negativereason_confidence	airline	\
1	NaN	0.0000	Virgin America	
3	Bad	0.7033	Virgin America	
4	Flight	1.0000	Virgin America	
5	Can't	0.6842	Virgin America	
6	Tell	0.0000	Virgin America	

	airline_sentiment	neg_gold	name	negativereason_gold	retweet_count	\
1	NaN	NaN	jnardino	NaN	0	
3	NaN	NaN	jnardino	NaN	0	
4	NaN	NaN	jnardino	NaN	0	
5	NaN	NaN	jnardino	NaN	0	
6	NaN	NaN	cjmcginnis	NaN	0	

	text	tweet_coord	\
1	@VirginAmerica plus you've added commercials t...	NaN	
3	@VirginAmerica it's really aggressive to blast...	NaN	
4	@VirginAmerica and it's a really big bad thing...	NaN	
5	@VirginAmerica seriously would pay \$30 a fligh...	NaN	
6	@VirginAmerica yes, nearly every time I fly VX...	NaN	

	tweet_created	tweet_location	user_timezone	\
1	2015-02-24	NaN	Pacific Time (US & Canada)	
3	2015-02-24	NaN	Pacific Time (US & Canada)	
4	2015-02-24	NaN	Pacific Time (US & Canada)	
5	2015-02-24	NaN	Pacific Time (US & Canada)	
6	2015-02-24	San Francisco CA	Pacific Time (US & Canada)	

	cleaned_tweet
1	you ve added commercials to the experience tacky

```

3          s really aggressive to blast obnoxious
4          enterta...
5          it s a really big bad thing about it
6          would pay a flight for seats that didn t have

```

```

[28]: df['cleaned_tweet'] = df['cleaned_tweet'].apply(text_process)
df.reset_index(drop=True, inplace = True)
df.head()

```

```

[28]:      index  tweet_id airline_sentiment airline_sentiment_confidence \
0         1  570301130888122368          1          0.3486
1         3  570301031407624196          0          1.0000
2         4  570300817074462722          0          1.0000
3         5  570300767074181121          0          1.0000
4         6  570300616901320704          1          0.6745

```

```

      negativereason negativereason_confidence      airline \
0          NaN          0.0000 Virgin America
1         Bad          0.7033 Virgin America
2        Flight          1.0000 Virgin America
3        Can't          0.6842 Virgin America
4         Tell          0.0000 Virgin America

```

```

      Can't
      airline_sentiment      name negativereason_gold retweet_count \
0          NaN      NaN      jnardino      NaN      0
1          NaN      NaN      jnardino      NaN      0
2          NaN      NaN      jnardino      NaN      0
3          NaN      NaN      jnardino      NaN      0
4          NaN      NaN      cjmccinnis      NaN      0

```

```

      text tweet_coord \
0  @VirginAmerica plus you've added commercials t...      NaN
1  @VirginAmerica it's really aggressive to blast...      NaN
2  @VirginAmerica and it's a really big bad thing...      NaN
3  @VirginAmerica seriously would pay $30 a fligh...      NaN
4  @VirginAmerica yes, nearly every time I fly VX...      NaN

```

```

      tweet_created tweet_location      user_timezone \
0  2015-02-24      NaN Pacific Time (US & Canada)
1  2015-02-24      NaN Pacific Time (US & Canada)
2  2015-02-24      NaN Pacific Time (US & Canada)
3  2015-02-24      NaN Pacific Time (US & Canada)
4  2015-02-24 San Francisco CA Pacific Time (US & Canada)

```

```

      cleaned_tweet
0      added
1      commercials
      experience tacky
      really aggressive
      blast obnoxious
      entertainmen...

```

2	really big bad thing
3	would pay flight seats
4	playing really bad thin...
	nearly every time fly vx
	ear worm go away

[29]: df['airline_sentiment'].unique()

[29]: array([1, 0])

2 Base SVM model with TF-IDF

[30]: # Creating object of TF-IDF vectorizer

```
vectorizer = TfidfVectorizer(use_idf=True, lowercase=True)
X_tf_idf = vectorizer.fit_transform(df.cleaned_tweet)
x_train, x_test, y_train, y_test = train_test_split(X_tf_idf,
df['airline_sentiment'], random_state=42)
```

[31]: SVM = svm.SVC(probability=True)

```
s_train_accuracy, s_test_accuracy, s_train_auc, s_test_auc =
check_scores(SVM, x_train, x_test, y_train, y_test)
```

Train confusion matrix is:

```
[[6824  31]
 [ 151 1649]]
```

Test confusion matrix is:

```
[[2291  32]
 [ 296 267]]
```

	precision	recall	f1-score	support
0	0.89	0.99	0.93	2323
1	0.89	0.47	0.62	563
accuracy			0.89	2886
macro avg	0.89	0.73	0.78	2886
weighted avg	0.89	0.89	0.87	2886

Train accuracy score: 0.9789716926632005


Test accuracy score: 0.8863478863478863

Train ROC-AUC score: 0.9969059080962801

Test ROC-AUC score: 0.9291699576938928

Are under Precision-Recall curve: 0.6194895591647333

Area under ROC-AUC: 0.8049981076642363



3 optimizing the hyperparameters with TF-IDF

```
[32]: parameters={
    "C":[0.1,1,10],
    "kernel":["linear", 'rbf', 'sigmoid'],
    "gamma":["scale", 'auto']
}

svm_optimal = grid_search(svm.SVC(probability=True), parameters,x_train,
    y_train)
```

Fitting 2 folds for each of 18 candidates, totalling 36 fits

[CV] END ...C=0.1, gamma=scale, kernel=linear; total time=	6.0s
[CV] END ...C=0.1, gamma=scale, kernel=linear; total time=	5.8s
[CV] END ...C=0.1, gamma=scale, kernel=rbf; total time=	9.6s
[CV] END ...C=0.1, gamma=scale, kernel=rbf; total time=	9.4s
[CV] END ...C=0.1, gamma=scale, kernel=sigmoid; total time=	5.9s
[CV] END ...C=0.1, gamma=scale, kernel=sigmoid; total time=	5.8s
[CV] END ...C=0.1, gamma=auto, kernel=linear; total time=	6.0s
[CV] END ...C=0.1, gamma=auto, kernel=linear; total time=	5.9s
[CV] END ...C=0.1, gamma=auto, kernel=rbf; total time=	4.7s
[CV] END ...C=0.1, gamma=auto, kernel=rbf; total time=	4.7s
[CV] END ...C=0.1, gamma=auto, kernel=sigmoid; total time=	4.4s

```

[CV] END ...C=0.1, gamma=auto, kernel=sigmoid; total time= 4.4s
[CV] END ...C=1, gamma=scale, kernel=linear; total time= 5.9s
[CV] END ...C=1, gamma=scale, kernel=linear; total time= 5.7s
[CV] END ...C=1, gamma=scale, kernel=rbf; total time= 10.6s
[CV] END ...C=1, gamma=scale, kernel=rbf; total time= 10.5s
[CV] END ...C=1, gamma=scale, kernel=sigmoid; total time= 5.6s
[CV] END ...C=1, gamma=scale, kernel=sigmoid; total time= 5.4s
[CV] END ...C=1, gamma=auto, kernel=linear; total time= 5.9s
[CV] END ...C=1, gamma=auto, kernel=linear; total time= 5.7s
[CV] END ...C=1, gamma=auto, kernel=rbf; total time= 4.9s
[CV] END ...C=1, gamma=auto, kernel=rbf; total time= 4.9s
[CV] END ...C=1, gamma=auto, kernel=sigmoid; total time= 4.5s
[CV] END ...C=1, gamma=auto, kernel=sigmoid; total time= 4.4s
[CV] END ...C=10, gamma=scale, kernel=linear; total time= 6.3s
[CV] END ...C=10, gamma=scale, kernel=linear; total time= 6.3s
[CV] END ...C=10, gamma=scale, kernel=rbf; total time= 11.6s
[CV] END ...C=10, gamma=scale, kernel=rbf; total time= 11.5s
[CV] END ...C=10, gamma=scale, kernel=sigmoid; total time= 8.8s
[CV] END ...C=10, gamma=scale, kernel=sigmoid; total time= 7.6s
[CV] END ...C=10, gamma=auto, kernel=linear; total time= 6.3s
[CV] END ...C=10, gamma=auto, kernel=linear; total time= 6.2s
[CV] END ...C=10, gamma=auto, kernel=rbf; total time= 6.2s
[CV] END ...C=10, gamma=auto, kernel=rbf; total time= 6.1s
[CV] END ...C=10, gamma=auto, kernel=sigmoid; total time= 5.3s
[CV] END ...C=10, gamma=auto, kernel=sigmoid; total time= 5.3s

```

Best parameters are:

```
{'C': 10, 'gamma': 'scale', 'kernel': 'rbf'}
```

```
[33]: so_train_accuracy, so_test_accuracy, so_train_auc, so_test_auc = _
      check_scores(svm_optimal, x_train, x_test, y_train, y_test)
```

Train confusion matrix is:

```
[[6829 26]
 [ 5 1795]]
```

Test confusion matrix is:

```
[[2272 51]
 [245 318]]
```

	precision	recall	f1-score	support
0	0.90	0.98	0.94	2323
1	0.86	0.56	0.68	563
accuracy			0.90	2886
macro avg	0.88	0.77	0.81	2886

weighted avg 0.89 0.90 0.89 2886

Train accuracy score: 0.996418255343732



Test accuracy score: 0.8974358974358975

Train ROC-AUC score: 0.9987310154793744

Test ROC-AUC score: 0.928759359834354

Are under Precision-Recall curve: 0.6824034334763949

Area under ROC-AUC: 0.80756377046123



4 Multinomial Naive Bayes

```
[34]: m_train_accuracy, m_test_accuracy, m_train_auc, m_test_auc = _  
      check_scores(MultinomialNB(),x_train, x_test, y_train, y_test)
```

Train confusion matrix is:

```
[[6853   2]  
 [1296 504]]
```

Test confusion matrix is:

```
[[2318   5]  
 [ 474  89]]
```

precision

recall f1-score

support

0	0.83	1.00	0.91	2323
1	0.95	0.16	0.27	563
accuracy			0.83	2886
macro avg	0.89	0.58	0.59	2886
weighted	0.85	0.83	0.78	2886
avg				

Train accuracy score: 0.8500288850375506



Test accuracy score: 0.834026334026334

Train ROC-AUC score: 0.9561105438041981

Test ROC-AUC score: 0.9013005324009118

Are under Precision-Recall curve: 0.2709284627092846

Area under ROC-AUC: 0.7374839454006489



5 Gaussian Naive Bayes

```
[35]: g_train_accuracy, g_test_accuracy, g_train_auc, \
      g_test_auc=check_scores(GaussianNB(),x_train.toarray(), x_test.toarray(), \
      y_train, y_test)
```

Train confusion matrix is:

[[5543 1312]

[0 1800]]

Test confusion matrix is:

[[1623 700]

[181 382]]

	precision	recall	f1-score	support
0	0.90	0.70	0.79	2323
1	0.35	0.68	0.46	563
accuracy			0.69	2886
macro avg	0.63	0.69	0.63	2886
weighted avg	0.79	0.69	0.72	2886

Train accuracy score: 0.8484113229347198



Test accuracy score: 0.6947331947331947

Train ROC-AUC score: 0.9043034281546316

Test ROC-AUC score: 0.688586755810495

Area under Precision-Recall curve: 0.4644376899696049

Area under ROC-AUC: 0.5471372315951626



It is interesting to see in Naive Bayes, we are getting linear relationship.

6 AdaBoost

```
[36]: a_train_accuracy, a_test_accuracy, a_train_auc, a_test_auc=check_scores(AdaBoostClassifier(),x_train,x_test, y_train, y_test)
```

Train confusion matrix is:

```
[[6655 200]
 [1012 788]]
```

Test confusion matrix is:

```
[[2251    72]
 [ 328 235]]
```

	precision	recall	f1-score	support
0	0.87	0.97	0.92	2323
1	0.77	0.42	0.54	563
accuracy			0.86	2886
macro avg	0.82	0.69	0.73	2886
weighted avg	0.85	0.86	0.84	2886

Train accuracy score: 0.8599653379549393


Test accuracy score: 0.8613998613998614

Train ROC-AUC score: 0.8689482940270687

Test ROC-AUC score: 0.8373803091947158

Are under Precision-Recall curve: 0.5402298850574713

Area under ROC-AUC: 0.6298876361965743



7 AdaBoost with hyperparameters

```
[37]: params = {'n_estimators': [10, 50, 100, 500],
               'learning_rate': [0.0001, 0.001, 0.01, 0.1, 1.0],
               'algorithm': ['SAMME', 'SAMME.R']}
```

```
ada_optimal_model = grid_search(AdaBoostClassifier(), params, x_train, y_train)
```

Fitting 2 folds for each of 40 candidates, totalling 80 fits

[CV] END algorithm=SAMME, learning_rate=0.0001, n_estimators=10; total time=0.4s

[CV] END algorithm=SAMME, learning_rate=0.0001, n_estimators=10; total time=0.4s

[CV] END algorithm=SAMME, learning_rate=0.0001, n_estimators=50; total time=1.8s

[CV] END algorithm=SAMME, learning_rate=0.0001, n_estimators=50; total time=1.7s

[CV] END algorithm=SAMME, learning_rate=0.0001, n_estimators=100; total time=3.5s

[CV] END algorithm=SAMME, learning_rate=0.0001, n_estimators=100; total time=3.5s

[CV] END algorithm=SAMME, learning_rate=0.0001, n_estimators=500; total time=17.6s

[CV] END algorithm=SAMME, learning_rate=0.0001, n_estimators=500; total time=17.5s

[CV] END algorithm=SAMME, learning_rate=0.001, n_estimators=10; total time=0.4s

[CV] END algorithm=SAMME, learning_rate=0.001, n_estimators=10; total time=0.4s

[CV] END algorithm=SAMME, learning_rate=0.001, n_estimators=50; total time=1.7s

[CV] END algorithm=SAMME, learning_rate=0.001, n_estimators=50; total time=1.7s

[CV] END algorithm=SAMME, learning_rate=0.001, n_estimators=100; total time=3.5s

[CV] END algorithm=SAMME, learning_rate=0.001, n_estimators=100; total time=3.5s

[CV] END algorithm=SAMME, learning_rate=0.001, n_estimators=500; total time=17.6s

[CV] END algorithm=SAMME, learning_rate=0.001, n_estimators=500; total time=17.4s

[CV] END algorithm=SAMME, learning_rate=0.01, n_estimators=10; total time=0.4s

[CV] END algorithm=SAMME, learning_rate=0.01, n_estimators=10; total time=0.4s

[CV] END algorithm=SAMME, learning_rate=0.01, n_estimators=50; total time=1.8s

[CV] END algorithm=SAMME, learning_rate=0.01, n_estimators=50; total time=1.7s

[CV] END algorithm=SAMME, learning_rate=0.01, n_estimators=100; total time=3.5s

[CV] END algorithm=SAMME, learning_rate=0.01, n_estimators=100; total time=3.5s

[CV] END algorithm=SAMME, learning_rate=0.01, n_estimators=500; total time=17.6s

[CV] END algorithm=SAMME, learning_rate=0.01, n_estimators=500; total time=17.5s

[CV] END algorithm=SAMME, learning_rate=0.1, n_estimators=10; total time=0.4s

[CV] END algorithm=SAMME, learning_rate=0.1, n_estimators=10; total time=0.4s

[CV] END algorithm=SAMME, learning_rate=0.1, n_estimators=50; total time=1.7s

[CV] END algorithm=SAMME, learning_rate=0.1, n_estimators=50; total time=1.8s

[CV] END algorithm=SAMME, learning_rate=0.1, n_estimators=100; total time=3.5s

[CV] END algorithm=SAMME, learning_rate=0.1, n_estimators=100; total time=3.5s

[CV] END algorithm=SAMME, learning_rate=0.1, n_estimators=500; total time=17.6s

[CV] END algorithm=SAMME, learning_rate=0.1, n_estimators=500; total time=17.5s

[CV] END algorithm=SAMME, learning_rate=1.0, n_estimators=10; total time=0.4s

[CV] END algorithm=SAMME, learning_rate=1.0, n_estimators=10; total time=0.4s

[CV] END algorithm=SAMME, learning_rate=1.0, n_estimators=50; total time=1.8s

[CV] END algorithm=SAMME, learning_rate=1.0, n_estimators=50; total time=1.7s

[CV] END algorithm=SAMME, learning_rate=1.0, n_estimators=100; total time=3.5s
[CV] END algorithm=SAMME, learning_rate=1.0, n_estimators=100; total time=3.5s
[CV] END algorithm=SAMME, learning_rate=1.0, n_estimators=500; total time=17.6s
[CV] END algorithm=SAMME, learning_rate=1.0, n_estimators=500; total time=17.5s
[CV] END algorithm=SAMME.R, learning_rate=0.0001, n_estimators=10; total time=0.4s
[CV] END algorithm=SAMME.R, learning_rate=0.0001, n_estimators=10; total time=0.4s
[CV] END algorithm=SAMME.R, learning_rate=0.0001, n_estimators=50; total time=1.8s
[CV] END algorithm=SAMME.R, learning_rate=0.0001, n_estimators=50; total time=1.8s
[CV] END algorithm=SAMME.R, learning_rate=0.0001, n_estimators=100; total time=3.6s
[CV] END algorithm=SAMME.R, learning_rate=0.0001, n_estimators=100; total time=3.6s
[CV] END algorithm=SAMME.R, learning_rate=0.0001, n_estimators=500; total time=17.9s
[CV] END algorithm=SAMME.R, learning_rate=0.0001, n_estimators=500; total time=17.9s
[CV] END algorithm=SAMME.R, learning_rate=0.001, n_estimators=10; total time=0.4s
[CV] END algorithm=SAMME.R, learning_rate=0.001, n_estimators=10; total time=0.4s
[CV] END algorithm=SAMME.R, learning_rate=0.001, n_estimators=50; total time=1.8s
[CV] END algorithm=SAMME.R, learning_rate=0.001, n_estimators=50; total time=1.8s
[CV] END algorithm=SAMME.R, learning_rate=0.001, n_estimators=100; total time=3.6s
[CV] END algorithm=SAMME.R, learning_rate=0.001, n_estimators=100; total time=3.6s
[CV] END algorithm=SAMME.R, learning_rate=0.001, n_estimators=500; total time=18.0s
[CV] END algorithm=SAMME.R, learning_rate=0.001, n_estimators=500; total time=17.9s
[CV] END algorithm=SAMME.R, learning_rate=0.01, n_estimators=10; total time=0.4s
[CV] END algorithm=SAMME.R, learning_rate=0.01, n_estimators=10; total time=0.4s
[CV] END algorithm=SAMME.R, learning_rate=0.01, n_estimators=50; total time=1.8s
[CV] END algorithm=SAMME.R, learning_rate=0.01, n_estimators=50; total time=1.8s

[CV] END algorithm=SAMME.R, learning_rate=0.01, n_estimators=100; total time=3.6s
 [CV] END algorithm=SAMME.R, learning_rate=0.01, n_estimators=100; total time=3.6s
 [CV] END algorithm=SAMME.R, learning_rate=0.01, n_estimators=500; total time=17.9s
 [CV] END algorithm=SAMME.R, learning_rate=0.01, n_estimators=500; total time=17.9s
 [CV] END algorithm=SAMME.R, learning_rate=0.1, n_estimators=10; total time=0.4s
 [CV] END algorithm=SAMME.R, learning_rate=0.1, n_estimators=10; total time=0.4s
 [CV] END algorithm=SAMME.R, learning_rate=0.1, n_estimators=50; total time=1.8s
 [CV] END algorithm=SAMME.R, learning_rate=0.1, n_estimators=50; total time=1.8s
 [CV] END algorithm=SAMME.R, learning_rate=0.1, n_estimators=100; total time=3.6s
 [CV] END algorithm=SAMME.R, learning_rate=0.1, n_estimators=100; total time=3.6s
 [CV] END algorithm=SAMME.R, learning_rate=0.1, n_estimators=500; total time=18.0s
 [CV] END algorithm=SAMME.R, learning_rate=0.1, n_estimators=500; total time=17.9s
 [CV] END algorithm=SAMME.R, learning_rate=1.0, n_estimators=10; total time=0.4s
 [CV] END algorithm=SAMME.R, learning_rate=1.0, n_estimators=10; total time=0.4s
 [CV] END algorithm=SAMME.R, learning_rate=1.0, n_estimators=50; total time=1.8s
 [CV] END algorithm=SAMME.R, learning_rate=1.0, n_estimators=50; total time=1.8s
 [CV] END algorithm=SAMME.R, learning_rate=1.0, n_estimators=100; total time=3.6s
 [CV] END algorithm=SAMME.R, learning_rate=1.0, n_estimators=100; total time=3.5s
 [CV] END algorithm=SAMME.R, learning_rate=1.0, n_estimators=500; total time=17.9s
 [CV] END algorithm=SAMME.R, learning_rate=1.0, n_estimators=500; total time=17.9s

Best parameters are:

```
{'algorithm': 'SAMME.R', 'learning_rate': 0.1, 'n_estimators': 500}
```

```
[38]: ao_train_accuracy, ao_test_accuracy, ao_train_auc, _  
      ao_test_auc=check_scores(ada_optimal_model,x_train,x_test,y_train,y_test)
```

Train confusion matrix is:

```
[[6761   94]
 [1043 757]]
```

Test confusion matrix is:

```
[[2278   45]
 [ 340 223]]
```

	precision	recall	f1-score	support
0	0.87	0.98	0.92	2323
1	0.83	0.40	0.54	563
accuracy			0.87	2886
macro avg	0.85	0.69	0.73	2886
weighted avg	0.86	0.87	0.85	2886

Train accuracy score: 0.868630849220104


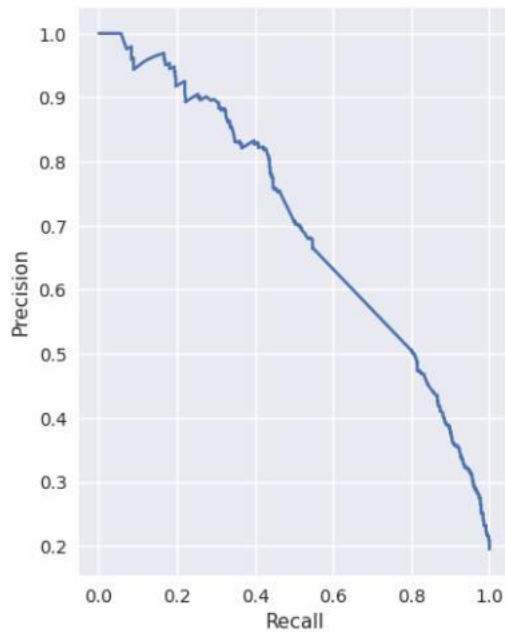
Test accuracy score: 0.8665973665973666

Train ROC-AUC score: 0.9218550936056407

Test ROC-AUC score: 0.8779339205061136

Are under Precision-Recall curve: 0.5367027677496993

Area under ROC-AUC: 0.7021179688268547



8 KNeighbors

```
[39]: from sklearn.neighbors import KNeighborsClassifier
      knn_train_accuracy, knn_test_accuracy, knn_train_auc, knn_test_auc=
      check_scores(KNeighborsClassifier().fit(x_train, y_train),
      x_train,x_test,y_train,y_test)
```

Train confusion matrix is:

```
[[ 661 6194]
 [   20 1780]]
```

Test confusion matrix is:

```
[[ 73 2250]
 [   6  557]]
```

	precision	recall	f1-score	support
0	0.92	0.03	0.06	2323
1	0.20	0.99	0.33	563
accuracy			0.22	2886
macro avg	0.56	0.51	0.20	2886
weighted avg	0.78	0.22	0.11	2886

Train accuracy score: 0.28203350664355864



Test accuracy score: 0.2182952182952183

Train ROC-AUC score: 0.9672222222222223

Test ROC-AUC score: 0.5688336344639173

Are under Precision-Recall curve: 0.3305637982195846

Area under ROC-AUC: 0.5761885418599553



9 Random Forest

```
[40]: r_train_accuracy, r_test_accuracy, r_train_auc, r_test_auc=
      check_scores(RandomForestClassifier(random_state=0).fit(x_train, y_train),
      x_train, x_test, y_train, y_test)
```

Train confusion matrix is:

```
[[6829 26]
 [ 5 1795]]
```

Test confusion matrix is:

```
[[2215 108]
 [ 238 325]]
```

	precision	recall	f1-score	support
0	0.90	0.95	0.93	2323
1	0.75	0.58	0.65	563
accuracy			0.88	2886
macro avg	0.83	0.77	0.79	2886
weighted avg	0.87	0.88	0.87	2886

Train accuracy score: 0.996418255343732



Test accuracy score: 0.8801108801108801

Train ROC-AUC score: 0.9982442661479861

Test ROC-AUC score: 0.8956867344777572

Area under Precision-Recall curve: 0.6526104417670683

Area under ROC-AUC: 0.7441899264879837



10 Decision Tree

```
[41]: from sklearn.tree import DecisionTreeClassifier
      dt_train_accuracy, dt_test_accuracy, dt_train_auc, dt_test_auc = \
          check_scores(DecisionTreeClassifier(random_state=0).fit(x_train, y_train), \
              x_train, x_test, y_train, y_test)
```

Train confusion matrix is:

```
[[6829 26]
 [ 5 1795]]
```

Test confusion matrix is:

```
[[2039 284]
 [ 199 364]]
```

	precision	recall	f1-score	support
0	0.91	0.88	0.89	2323

	1	0.56	0.65	0.60	563
accuracy				0.83	2886
macro avg		0.74	0.76	0.75	2886
weighted		0.84	0.83	0.84	2886
avg					

Train accuracy score: 0.996418255343732


Test accuracy score: 0.8326403326403327

Train ROC-AUC score: 0.9998750303914418

Test ROC-AUC score: 0.7627015809929127

Area under Precision-Recall curve: 0.6011560693641619

Area under ROC-AUC: 0.6116197326239008



11 Neural Network

```
[42]: from sklearn.neural_network import MLPClassifier
      mlp_train_accuracy, mlp_test_accuracy, mlp_train_auc, mlp_test_auc = \
      check_scores(MLPClassifier(max_iter=500, random_state=42).fit(x_train, \
      y_train), x_train, x_test, y_train, y_test)
```

Train confusion matrix is:

```
[[6825 30]
```

```
 [ 5 1795]]
```

Test confusion matrix is:

[[2125 198]

[178 385]]

	precision	recall	f1-score	support
0	0.92	0.91	0.92	2323
1	0.66	0.68	0.67	563
accuracy			0.87	2886
macro avg	0.79	0.80	0.80	2886
weighted avg	0.87	0.87	0.87	2886

Train accuracy score: 0.9959560947429231



Test accuracy score: 0.8697158697158697

Train ROC-AUC score: 0.9998520139395413

Test ROC-AUC score: 0.9014878628954872

Are under Precision-Recall curve: 0.6719022687609075

Area under ROC-AUC: 0.724167033127511



12 LSTM

```
[43]: corpus = [df['cleaned_tweet'][i] for i in range(len(df))]  
  
voc_size=5000  
  
onehot_=[one_hot(words,voc_size) for words in corpus]  
  
max_sent_length=max([len(i) for i in corpus])  
  
embedded_docs=pad_sequences(onehot_,padding='pre',maxlen=max_sent_length)  
  
embedding_vector_features=40  
model=Sequential()  
model.  
    add(Embedding(voc_size,embedding_vector_features,input_length=max_sent_length))  
    model.add(Dropout(0.3))  
    model.add(LSTM(100))  
model.add(Dropout(0.3))  
model.add(Dense(1,activation='sigmoid'))  
model.compile(loss='binary_crossentropy',optimizer='adam',metrics=['accuracy'])  
  
X_final=np.array(embedded_docs)  
y_final=np.array(df['airline_sentiment'])  
X_final.shape,y_final.shape
```

[43]: ((11541, 124), (11541,))

```
[44]: X_train, X_test, y_train, y_test = train_test_split(X_final, y_final,   
    test_size=0.33, random_state=42)  
model.  
    fit(X_train,y_train,validation_data=(X_test,y_test),epochs=10,batch_size=64)
```

Epoch 1/10

121/121 [=====] - 26s 192ms/step - loss: 0.4484 -
accuracy: 0.8100 - val_loss: 0.3276 - val_accuracy: 0.8645

Epoch 2/10

121/121 [=====] - 21s 175ms/step - loss: 0.2641 -
accuracy: 0.8927 - val_loss: 0.2858 - val_accuracy: 0.8777

Epoch 3/10

121/121 [=====] - 21s 175ms/step - loss: 0.1913 -
accuracy: 0.9245 - val_loss: 0.2932 - val_accuracy: 0.8777

Epoch 4/10

121/121 [=====] - 22s 178ms/step - loss: 0.1558 -
accuracy: 0.9419 - val_loss: 0.3312 - val_accuracy: 0.8719

Epoch 5/10

121/121 [=====] - 21s 176ms/step - loss: 0.1291 -

```

accuracy: 0.9531 - val_loss: 0.3519 - val_accuracy: 0.8727
Epoch 6/10
121/121 [=====] - 21s 173ms/step - loss: 0.1073 -
accuracy: 0.9613 - val_loss: 0.4079 - val_accuracy: 0.8819
Epoch 7/10
121/121 [=====] - 21s 176ms/step - loss: 0.0920 -
accuracy: 0.9659 - val_loss: 0.4288 - val_accuracy: 0.8784
Epoch 8/10
121/121 [=====] - 21s 176ms/step - loss: 0.0791 -
accuracy: 0.9718 - val_loss: 0.4780 - val_accuracy: 0.8748
Epoch 9/10
121/121 [=====] - 21s 175ms/step - loss: 0.0669 -
accuracy: 0.9766 - val_loss: 0.5347 - val_accuracy: 0.8669
Epoch 10/10
121/121 [=====] - 21s 174ms/step - loss: 0.0622 -
accuracy: 0.9784 - val_loss: 0.5220 - val_accuracy: 0.8677

```

[44]: <keras.callbacks.History at 0x731470a4bdf0>

```

[45]: y_test_pred=np.argmax(model.predict(X_test),axis=1)
y_train_pred=np.argmax(model.predict(X_train),axis=1)

```

```

120/120 [=====] - 4s 26ms/step
242/242 [=====] - 7s 27ms/step

```

```

[46]: test_acc_lstm = accuracy_score(y_test,y_test_pred)
train_acc_lstm = accuracy_score(y_train,y_train_pred)
test_roc_lstm = roc_auc_score(y_test,y_test_pred)
train_roc_lstm = roc_auc_score(y_train,y_train_pred)

```

```

[47]: data = [('Random Forest', r_train_accuracy, r_test_accuracy, r_train_auc,
r_test_auc),
('MultinomialNB',m_train_accuracy, m_test_accuracy, m_train_auc, m_test_auc ),
('KNeighbors',knn_train_accuracy, knn_test_accuracy, knn_train_auc,
knn_test_auc ),
('AdaBoost',a_train_accuracy, a_test_accuracy, a_train_auc, a_test_auc ),
('AdaBoost Optimized',ao_train_accuracy, ao_test_accuracy, ao_train_auc,
ao_test_auc),
('Decision Tree ',dt_train_accuracy, dt_test_accuracy, dt_train_auc,
dt_test_auc),
('Gaussian Naive Bayes',g_train_accuracy, g_test_accuracy, g_train_auc,
g_test_auc),
('SVM', s_train_accuracy, s_test_accuracy, s_train_auc, s_test_auc),
('SVM Optimized', so_train_accuracy, so_test_accuracy, so_train_auc,
so_test_auc),
('Neural Network',mlp_train_accuracy, mlp_test_accuracy, mlp_train_auc,
mlp_test_auc),

```

```

('LSTM',train_acc_lstm, test_acc_lstm, train_roc_lstm, test_roc_lstm)])

Scores_ =pd.DataFrame(data = data, columns=['Model Name','Train Accuracy', '_
        'Test Accuracy', 'Train ROC', 'Test ROC'])
        ↳ Scores_.set_index('Model Name', inplace = True)

Scores_

```

[47]:

	Train Accuracy	Test Accuracy	Train ROC	Test ROC
Model Name				
Random Forest	0.996418	0.880111	0.998244	0.895687
MultinomialNB	0.850029	0.834026	0.956111	0.901301
KNeighbors	0.282034	0.218295	0.967222	0.568834
AdaBoost	0.859965	0.861400	0.868948	0.837380
AdaBoost Optimized	0.868631	0.866597	0.921855	0.877934
Decision Tree	0.996418	0.832640	0.999875	0.762702
Gaussian Naive Bayes	0.848411	0.694733	0.904303	0.688587
SVM	0.978972	0.886348	0.996906	0.929170
SVM Optimized	0.996418	0.897436	0.998731	0.928759
Neural Network	0.995956	0.869716	0.999852	0.901488
LSTM	0.794232	0.797322	0.500000	0.500000

Conclusion

1 Improved Analysis Accuracy

The developed sentiment analysis model significantly improves accuracy compared to traditional manual analysis.

2 Efficient Solution

The model is efficient and scalable, reducing resource requirements and improving decision-making speed.

3 Further Improvements Possible

Although the model offers significant improvements, further tweaks and improvements are possible to achieve even better accuracy and efficiency.